

# Intermediate SQL

- **Slideds based on:**

- Database System Concepts

- Avi Silberschatz, Henry F. Korth, S. Sudarshan

- <http://www.db-book.com/>

# Intermediate SQL

- **Join Expressions**
- **Views**
- **Transactions**
- **Integrity Constraints**
- **SQL Data Types and Schemas**
- **Authorization**

# Joined Relations

- **Join operations** take two relations and return as a result another relation.
- A join operation is a Cartesian product which requires that tuples in the two relations match (under some condition). It also specifies the attributes that are present in the result of the join
- The join operations are typically used as subquery expressions in the from clause

# Join operations – Example

- Relation *course*

<i>course_id</i>	<i>title</i>	<i>dept_name</i>	<i>credits</i>
BIO-301	Genetics	Biology	4
CS-190	Game Design	Comp. Sci.	4
CS-315	Robotics	Comp. Sci.	3

- Relation *prereq*

<i>course_id</i>	<i>prereq_id</i>
BIO-301	BIO-101
CS-190	CS-101
CS-347	CS-101

- Observe that

prereq information is missing for CS-315 and  
course information is missing for CS-437

# Outer Join

- An extension of the join operation that avoids loss of information.
- Computes the join and then adds tuples from one relation that does not match tuples in the other relation to the result of the join.
- Uses *null* values.

# Left Outer Join

■ *course* **natural left outer join** *prereq*

<i>course_id</i>	<i>title</i>	<i>dept_name</i>	<i>credits</i>	<i>prereq_id</i>
BIO-301	Genetics	Biology	4	BIO-101
CS-190	Game Design	Comp. Sci.	4	CS-101
CS-315	Robotics	Comp. Sci.	3	<i>null</i>

# Right Outer Join

■ *course* **natural right outer join** *prereq*

<i>course_id</i>	<i>title</i>	<i>dept_name</i>	<i>credits</i>	<i>prereq_id</i>
BIO-301	Genetics	Biology	4	BIO-101
CS-190	Game Design	Comp. Sci.	4	CS-101
CS-347	<i>null</i>	<i>null</i>	<i>null</i>	CS-101



# Joined Relations

- **Join operations** take two relations and return as a result another relation.
- These additional operations are typically used as subquery expressions in the from clause
- **Join condition** – defines which tuples in the two relations match, and what attributes are present in the result of the join.
- **Join type** – defines how tuples in each relation that do not match any tuple in the other relation (based on the join condition) are treated.

<i>Join types</i>	<i>Join Conditions</i>
inner join left outer join right outer join full outer join	natural on <predicate> using ( $A_1, A_1, \dots, A_n$ )

# Full Outer Join

■ *course* **natural full outer join** *prereq*

<i>course_id</i>	<i>title</i>	<i>dept_name</i>	<i>credits</i>	<i>prereq_id</i>
BIO-301	Genetics	Biology	4	BIO-101
CS-190	Game Design	Comp. Sci.	4	CS-101
CS-315	Robotics	Comp. Sci.	3	<i>null</i>
CS-347	<i>null</i>	<i>null</i>	<i>null</i>	CS-101

# Joined Relations – Examples

- *course* inner join *prereq* on  
*course.course\_id* = *prereq.course\_id*

<i>course_id</i>	<i>title</i>	<i>dept_name</i>	<i>credits</i>	<i>prereq_id</i>	<i>course_id</i>
BIO-301	Genetics	Biology	4	BIO-101	BIO-301
CS-190	Game Design	Comp. Sci.	4	CS-101	CS-190

- What is the difference between the above, and a natural join?
- *course* **left outer join** *prereq* on  
*course.course\_id* = *prereq.course\_id*

<i>course_id</i>	<i>title</i>	<i>dept_name</i>	<i>credits</i>	<i>prereq_id</i>	<i>course_id</i>
BIO-301	Genetics	Biology	4	BIO-101	BIO-301
CS-190	Game Design	Comp. Sci.	4	CS-101	CS-190
CS-315	Robotics	Comp. Sci.	3	<i>null</i>	<i>null</i>

# Joined Relations – Examples

- *course* **natural right outer join** *prereq*

<i>course_id</i>	<i>title</i>	<i>dept_name</i>	<i>credits</i>	<i>prereq_id</i>
BIO-301	Genetics	Biology	4	BIO-101
CS-190	Game Design	Comp. Sci.	4	CS-101
CS-347	<i>null</i>	<i>null</i>	<i>null</i>	CS-101

- *course* **full outer join** *prereq* **using** (*course\_id*)

<i>course_id</i>	<i>title</i>	<i>dept_name</i>	<i>credits</i>	<i>prereq_id</i>
BIO-301	Genetics	Biology	4	BIO-101
CS-190	Game Design	Comp. Sci.	4	CS-101
CS-315	Robotics	Comp. Sci.	3	<i>null</i>
CS-347	<i>null</i>	<i>null</i>	<i>null</i>	CS-101

# Views

- In some cases, it is not desirable for all users to see the entire logical model (that is, all the actual relations stored in the database.)
- Consider a person who needs to know an instructors name and department, but not the salary. This person should see a relation described, in SQL, by

```
select ID, name, dept_name  
from instructor
```

- A **view** provides a mechanism to hide certain data from the view of certain users.
- Any relation that is not of the conceptual model but is made visible to a user as a “virtual relation” is called a **view**.

# View Definition

- A view is defined using the create view statement which has the form

**create view  $v$  as < query expression >**

where <query expression> is any legal SQL expression. The view name is represented by  $v$ .

- Once a view is defined, the view name can be used to refer to the virtual relation that the view generates.
- View definition is not the same as creating a new relation by evaluating the query expression
  - Rather, a view definition causes the saving of an expression; the expression is substituted into queries using the view.

# Example Views

- A view of instructors without their salary  
create view *faculty* as  
select *ID*, *name*, *dept\_name*  
from *instructor*
- Find all instructors in the Biology department  
select *name*  
from *faculty*  
where *dept\_name* = 'Biology'
- Create a view of department salary totals  
create view *departments\_total\_salary*(*dept\_name*, *total\_salary*) as  
select *dept\_name*, sum (*salary*)  
from *instructor*  
group by *dept\_name*;

# Views Defined Using Other Views

- create view *physics\_fall\_2009* as  
select *course.course\_id*, *sec\_id*, *building*, *room\_number*  
from *course*, *section*  
where *course.course\_id* = *section.course\_id*  
and *course.dept\_name* = 'Physics'  
and *section.semester* = 'Fall'  
and *section.year* = '2009';
- create view *physics\_fall\_2009\_watson* as  
select *course\_id*, *room\_number*  
from *physics\_fall\_2009*  
where *building* = 'Watson';



# View Expansion

- Expand use of a view in a query/another view

```
create view physics_fall_2009_watson as  
(select course_id, room_number  
from (select course.course_id, building, room_number  
      from course, section  
      where course.course_id = section.course_id  
           and course.dept_name = 'Physics'  
           and section.semester = 'Fall'  
           and section.year = '2009')  
where building= 'Watson';
```

# Views Defined Using Other Views

- One view may be used in the expression defining another view
- A view relation  $v_1$  is said to **depend directly** on a view relation  $v_2$  if  $v_2$  is used in the expression defining  $v_1$
- A view relation  $v_1$  is said to **depend on** view relation  $v_2$  if either  $v_1$  depends directly to  $v_2$  or there is a path of dependencies from  $v_1$  to  $v_2$
- A view relation  $v$  is said to be **recursive** if it depends on itself.

# View Expansion

- A way to define the meaning of views defined in terms of other views.
- Let view  $v_1$  be defined by an expression  $e_1$  that may itself contain uses of view relations.
- View expansion of an expression repeats the following replacement step:
  - repeat
    - Find any view relation  $v_i$  in  $e_1$
    - Replace the view relation  $v_i$  by the expression defining  $v_i$
  - until no more view relations are present in  $e_1$
- As long as the view definitions are not recursive, this loop will terminate

# Update of a View

- Add a new tuple to *faculty* view which we defined earlier  
insert into *faculty* values ('30765', 'Green', 'Music');  
This insertion must be represented by the insertion of the tuple  
( '30765', 'Green', 'Music', null)  
into the *instructor* relation

# Some Updates cannot be Translated Uniquely

- create view *instructor\_info* as  
select *ID, name, building*  
from *instructor, department*  
where *instructor.dept\_name= department.dept\_name*;
- insert into *instructor\_info* values ('69987', 'White', 'Taylor');
  - which department, if multiple departments in Taylor?
  - what if no department is in Taylor?
- **Most SQL implementations allow updates only on simple views**
  - The **from** clause has only one database relation.
  - The **select** clause contains only attribute names of the relation, and does not have any expressions, aggregates, or **distinct** specification.
  - Any attribute not listed in the **select** clause can be set to null
  - The query does not have a **group** by or **having** clause.

# And Some Not at All

- create view *history\_instructors* as  
select \*  
from *instructor*  
where *dept\_name*= 'History';
- What happens if we insert ('25566', 'Brown', 'Biology', 100000) into *history\_instructors*?

# Transactions

- **Unit of work**
- **Atomic transaction**
  - either fully executed or rolled back as if it never occurred
- **Isolation from concurrent transactions**
- **Transactions begin implicitly**
  - Ended by **commit work** or **rollback work**
- **But default on most databases: each SQL statement commits automatically**
  - Can turn off auto commit for a session (e.g. using API)
  - In SQL:1999, can use: **begin atomic .... end**
    - Not supported on most databases

# Integrity Constraints

- **Integrity constraints guard against accidental damage to the database, by ensuring that authorized changes to the database do not result in a loss of data consistency.**
  - A checking account must have a balance greater than \$10,000.00
  - A salary of a bank employee must be at least \$4.00 an hour
  - A customer must have a (non-null) phone number



# Integrity Constraints on a Single Relation

- not null
- primary key
- unique
- check (P), where P is a predicate

# Not Null and Unique Constraints

- **not null**
  - Declare *name* and *budget* to be **not null**  
***name* varchar(20) not null**  
***budget* numeric(12,2) not null**
- **unique (  $A_1, A_2, \dots, A_m$  )**
  - The unique specification states that the attributes  $A_1, A_2, \dots, A_m$  form a candidate key.
  - Candidate keys are permitted to be null (in contrast to primary keys).

# The check clause

- **check (P)**  
where P is a predicate

Example: ensure that semester is one of fall, winter, spring or summer:

```
create table section (  
    course_id varchar (8),  
    sec_id varchar (8),  
    semester varchar (6),  
    year numeric (4,0),  
    building varchar (15),  
    room_number varchar (7),  
    time slot id varchar (4),  
    primary key (course_id, sec_id, semester, year),  
    check (semester in ('Fall', 'Winter', 'Spring', 'Summer'))  
);
```

# Referential Integrity

- Ensures that a value that appears in one relation for a given set of attributes also appears for a certain set of attributes in another relation.
  - Example: If “Biology” is a department name appearing in one of the tuples in the *instructor* relation, then there exists a tuple in the *department* relation for “Biology”.
- Let **A** be a set of attributes. Let **R** and **S** be two relations that contain attributes **A** and where **A** is the primary key of **S**. **A** is said to be a **foreign key** of **R** if for any values of **A** appearing in **R** these values also appear in **S**.

# Cascading Actions in Referential Integrity

- **create table *course* (  
    *course\_id* char(5) primary key,  
    *title* varchar(20),  
    *dept\_name* varchar(20) references *department*  
)**
- **create table *course* (  
    ...  
    *dept\_name* varchar(20),  
    foreign key (*dept\_name*) references *department*  
        on delete cascade  
        on update cascade,  
    ...  
)**
- **alternative actions to cascade: set null, set default**

# Integrity Constraint Violation During Transactions

- E.g.

```
create table person (  
    ID char(10),  
    name char(40),  
    mother char(10),  
    father char(10),  
    primary key ID,  
    foreign key father references person,  
    foreign key mother references person)
```

- **How to insert a tuple without causing constraint violation ?**
  - insert father and mother of a person before inserting person
  - OR, set father and mother to null initially, update after inserting all persons (not possible if father and mother attributes declared to be **not null**)
  - OR defer constraint checking (next slide)

# Complex Check Clauses

- **check** (*time\_slot\_id* in (select *time\_slot\_id* from *time\_slot*))
  - why not use a foreign key here?
- **Every section has at least one instructor teaching the section.**
  - how to write this?
- **Unfortunately: subquery in check clause not supported by pretty much any database**
  - Alternative: triggers (later)
- **create assertion <assertion-name> check <predicate>;**
  - Also not supported by anyone

# Built-in Data Types in SQL

- **date**: Dates, containing a (4 digit) year, month and date
  - Example: **date** '2005-7-27'
- **time**: Time of day, in hours, minutes and seconds.
  - Example: **time** '09:00:30'      **time** '09:00:30.75'
- **timestamp**: date plus time of day
  - Example: **timestamp** '2005-7-27 09:00:30.75'
- **interval**: period of time
  - Example: **interval** '1' day
  - Subtracting a date/time/timestamp value from another gives an interval value
  - Interval values can be added to date/time/timestamp values



# Index Creation

- create table *student*  
(*ID* varchar (5),  
*name* varchar (20) not null,  
*dept\_name* varchar (20),  
*tot\_cred* numeric (3,0) default 0,  
primary key (*ID*))
- create index *studentID\_index* on *student*(*ID*)
- Indices are data structures used to speed up access to records with specified values for index attributes
  - e.g. **select** \*  
    **from** *student*  
    **where** *ID* = '12345'

can be executed by using the index to find the required record, without looking at all records of *student*

*More on indices in Chapter 11*

# User-Defined Types

- **create type** construct in SQL creates user-defined type

**create type** *Dollars* as numeric (12,2) final

- **create table** *department*  
(*dept\_name* **varchar** (20),  
*building* **varchar** (15),  
*budget* *Dollars*);

# Domains

- **create domain** construct in SQL-92 creates user-defined domain types

```
create domain person_name char(20) not null
```

- Types and domains are similar. Domains can have constraints, such as not null, specified on them.
- **create domain** *degree\_level* **varchar**(10)  
**constraint** *degree\_level\_test*  
**check** (value in ('Bachelors', 'Masters', 'Doctorate'));

# Large-Object Types

- Large objects (photos, videos, CAD files, etc.) are stored as a *large object*:
  - **blob**: binary large object -- object is a large collection of uninterpreted binary data (whose interpretation is left to an application outside of the database system)
  - **clob**: character large object -- object is a large collection of character data
  - When a query returns a large object, a pointer is returned rather than the large object itself.

# Authorization

Forms of authorization on parts of the database:

- **Read** - allows reading, but not modification of data.
- **Insert** - allows insertion of new data, but not modification of existing data.
- **Update** - allows modification, but not deletion of data.
- **Delete** - allows deletion of data.

Forms of authorization to modify the database schema

- **Index** - allows creation and deletion of indices.
- **Resources** - allows creation of new relations.
- **Alteration** - allows addition or deletion of attributes in a relation.
- **Drop** - allows deletion of relations.

# Authorization Specification in SQL

- The **grant** statement is used to confer authorization  
    **grant <privilege list>**  
    **on <relation name or view name> to <user list>**
- **<user list> is:**
  - a user-id
  - **public**, which allows all valid users the privilege granted
  - A role (more on this later)
- **Granting a privilege on a view does not imply granting any privileges on the underlying relations.**
- **The grantor of the privilege must already hold the privilege on the specified item (or be the database administrator).**

# Privileges in SQL

- **select**: allows read access to relation, or the ability to query using the view
  - Example: grant users  $U_1$ ,  $U_2$ , and  $U_3$  **select** authorization on the *instructor* relation:

grant select on *instructor* to  $U_1$ ,  $U_2$ ,  $U_3$

- **insert**: the ability to insert tuples
- **update**: the ability to update using the SQL update statement
- **delete**: the ability to delete tuples.
- **all privileges**: used as a short form for all the allowable privileges

# Revoking Authorization in SQL

- The **revoke** statement is used to revoke authorization.  
    **revoke** <privilege list>  
    **on** <relation name or view name> **from** <user list>
- **Example:**  
    **revoke select on** *branch* **from**  $U_1, U_2, U_3$
- <privilege-list> may be **all** to revoke all privileges the revokee may hold.
- If <revokee-list> includes **public**, all users lose the privilege except those granted it explicitly.
- If the same privilege was granted twice to the same user by different grantees, the user may retain the privilege after the revocation.
- All privileges that depend on the privilege being revoked are also revoked.



# Roles

- **create role** instructor;
- **grant** *instructor* **to** Amit;
- **Privileges can be granted to roles:**
  - **grant** **select on** *takes* **to** *instructor*;
- **Roles can be granted to users, as well as to other roles**
  - **create role** *teaching\_assistant*
  - **grant** *teaching\_assistant* **to** *instructor*;
    - *Instructor* inherits all privileges of *teaching\_assistant*
- **Chain of roles**
  - **create role** *dean*;
  - **grant** *instructor* **to** *dean*;
  - **grant** *dean* **to** Satoshi;

# Authorization on Views

- create view *geo\_instructor* as  
(select \*  
from *instructor*  
where *dept\_name* = 'Geology');
- grant select on *geo\_instructor* to *geo\_staff*
- Suppose that a *geo\_staff* member issues
  - select \*  
from *geo\_instructor*;
- What if
  - *geo\_staff* does not have permissions on *instructor*?
  - creator of view did not have some permissions on *instructor*?

# Other Authorization Features

- **references privilege to create foreign key**
  - **grant reference** (*dept\_name*) **on** *department* **to** Mariano;
  - why is this required?
- **transfer of privileges**
  - **grant select on** *department* **to** Amit **with grant option**;
  - **revoke select on** *department* **from** Amit, Satoshi **cascade**;
  - **revoke select on** *department* **from** Amit, Satoshi **restrict**;
- **Etc. read Section 4.6 for more details we have omitted here.**

